

What is Spark Shell?

- **Spark Shell** is an **interactive REPL (Read-Eval-Print Loop)** for Apache Spark.
- It uses **Scala** by default and lets you test Spark commands instantly.
- It connects automatically to a **SparkContext** (sc) and **SparkSession** (spark in Spark 2.0+).

RDD Spark Tutorial

RDD ([Resilient Distributed Dataset](#)) is a fundamental data structure of Spark and it is the primary data abstraction in Apache Spark and the Spark Core. **RDDs are fault-tolerant, immutable distributed collections of objects, which means once you create an RDD you cannot change it.** Each dataset in RDD is divided into logical partitions, which can be computed on different nodes of the cluster.

RDD creation

RDDs are created primarily in two different ways, first parallelizing an existing collection and secondly referencing a dataset in an external storage system (HDFS, HDFS, S3 and many more).

sparkContext.parallelize()

sparkContext.parallelize is used to parallelize an existing collection in your driver program. This is a basic method to create RDD.

```
//Create RDD from parallelize
```

```
val dataSeq = Seq(("Java", 20000), ("Python", 100000), ("Scala", 3000))
```

```
val rdd=spark.sparkContext.parallelize(dataSeq)
```

sparkContext.textFile()

Using textFile() method we can read a text (.txt) file from many sources like HDFS, S#, Azure, local e.t.c into RDD.

```
//Create RDD from external Data source
```

```
val rdd2 = spark.sparkContext.textFile("/path/textFile.txt")
```

RDD Operations

On Spark RDD, you can perform **two kinds of operations**.

RDD Transformations

[Spark RDD Transformations](#) are lazy operations meaning they don't execute until you call an action on RDD. Since RDDs are immutable, When you run a transformation(for example map()), instead of updating a current RDD, it returns a new RDD.

Some transformations on RDDs are flatMap(), map(), reduceByKey(), filter(), sortByKey() and all these return a new RDD instead of updating the current.

RDD Actions

[RDD Action operation](#) returns the values from an RDD to a driver node. In other words, any RDD function that returns non RDD[T] is considered as an action. RDD operations trigger the computation and return RDD in a List to the driver program.

Some actions on RDDs are count(), collect(), first(), max(), reduce() and more.

RDD Lineage

What is Lineage?

- **RDD lineage** refers to the **sequence of transformations** that created an RDD.
- It forms a **DAG (Directed Acyclic Graph)**.
- If a partition of an RDD is lost, Spark **recomputes it** using its lineage.

Example of Lineage:

```
val rdd1 = sc.textFile("data.txt")           // 1. Load
val rdd2 = rdd1.flatMap(line => line.split(" ")) // 2. Split words
val rdd3 = rdd2.map(word => (word, 1))         // 3. Map to (word, 1)
val rdd4 = rdd3.reduceByKey(_ + _)            // 4. Count
```

```
val result = rdd4.collect()
```

```
// 5. Trigger action
```

Here's the **lineage DAG**:

```
textFile("data.txt")
```

↓

```
flatMap
```

↓

```
map
```

↓

```
reduceByKey
```

↓

```
collect()
```

What is Partitioning?

Partitioning refers to **how data is split across the cluster**. Spark processes data **partition-by-partition** in parallel.

Why Partitioning Matters:

- Controls **data locality**.
- Reduces **data shuffling** in operations like join, groupByKey, etc.
- Enables **parallel processing** for faster computation.

Persistence and Caching in Spark

What is Persistence?

Persistence is used to **store intermediate RDD/DataFrame results** in memory or disk so they don't get recomputed.

Why Persist?

- Avoid recomputation of lineage.
- Improve performance in iterative algorithms (e.g., ML, GraphX).
- Used when same RDD is accessed **multiple times**.

What is a DAG in Spark?

A **DAG (Directed Acyclic Graph)** in Spark represents:

A **logical execution plan** of **RDD transformations**, where each **node is an RDD** and **edges are operations** (like map, filter, etc.).

DAG Characteristics:

- **Directed**: Data flows in one direction (from source to result).
- **Acyclic**: No loops/cycles — once a transformation is applied, it moves forward.
- **Graph**: Shows dependencies among RDDs.

How It Works:

1. **User writes code with transformations** (e.g. map, filter, reduceByKey)
2. Spark **doesn't execute immediately** — it builds a **DAG** in memory.
3. When an **action** is called (e.g. collect(), count()), Spark:
 - Submits the DAG to the **DAG Scheduler**
 - Divides it into **stages**
 - Further divides stages into **tasks**
 - Executes tasks in **parallel** using **Task Scheduler**

Example DAG

```
val rdd1 = sc.textFile("data.txt")
val rdd2 = rdd1.flatMap(_.split(" "))
val rdd3 = rdd2.map(word => (word, 1))
```

```
val rdd4 = rdd3.reduceByKey(_ + _)
val rdd5 = rdd4.mapValues(_ * 2)
val result = rdd5.collect()
```

DAG Visualization

textFile("data.txt")

↓

flatMap

↓

map (word, 1)

↓

reduceByKey

↓

mapValues(*2)

↓

collect()

DAG vs Lineage

Concept Description

DAG Execution plan (logical graph)

Lineage History of how an RDD was created

Lineage is **used for fault recovery**, DAG is used for **execution planning**.

DAG Execution Flow

1. Logical Plan (DAG creation)

- Built when you chain transformations.

2. DAG Scheduler

- Splits the DAG into **stages** based on wide/narrow dependencies.

- **Narrow dependency:** No shuffle (e.g. map, filter)
- **Wide dependency:** Shuffle needed (e.g. reduceByKey)

3. Stage Division

- Each **stage** contains a **set of tasks** that can be executed in parallel.

4. Task Scheduler

- Assigns tasks to executors.

Component Role

DAG	Represents the complete job graph
Stage	A set of transformations that can be pipelined
Task	A single unit of execution (on a partition)
Job	Created for each action

RDD Programming Patterns in Scala

1. Transformation Patterns

Transformations are **lazy operations** that define a new RDD from the existing one.

***map**- Apply a function to each elements

```
val rdd = sc.parallelize(1 to 5)
```

```
val squared = rdd.map(x => x * x)
```

***flatMap**-Split elements into multiple items.

```
val lines = sc.parallelize(Seq("hello world", "spark scala"))
```

```
val words = lines.flatMap(_.split(" "))
```

***filter** -Filter elements based on a predicate.

```
val even = rdd.filter(_ % 2 == 0)
```

***distinct**-Removes duplicates

```
val distinctRDD = sc.parallelize(Seq(1, 2, 2, 3)).distinct()
```

* **union, intersection, subtract**- Set operations on RDDs.

```
val rdd1 = sc.parallelize(1 to 5)
```

```
val rdd2 = sc.parallelize(3 to 7)
```

```
val union = rdd1.union(rdd2)
```

```
val intersection = rdd1.intersection(rdd2)
```

```
val diff = rdd1.subtract(rdd2)
```

2. Key-Value (Pair RDD) Patterns

- **Key-based transformations and aggregations.**

* **mapToPair / map to tuple**

```
val pairs = sc.parallelize(Seq("apple", "banana"))
```

```
.map(word => (word.length, word))
```

* **reduceByKey** - Aggregates values with the same key.

```
val data = sc.parallelize(Seq(("math", 90), ("math", 80), ("eng", 85)))
```

```
val totals = data.reduceByKey(_ + _)
```

* **groupByKey** - Groups all values with the same key. (Less efficient)

```
val grouped = data.groupByKey()
```

* **sortByKey**

```
val sorted = data.sortByKey()
```

Joins in Spark (Scala)

Joins are operations to combine datasets based on a key.

1. RDD Joins (Key-Value RDDs)

- ♦ **join** – Inner Join-Returns only matching keys from both RDDs.

```
val rdd1 = sc.parallelize(Seq((1, "Alice"), (2, "Bob")))
val rdd2 = sc.parallelize(Seq((1, "Math"), (3, "Science")))

val joined = rdd1.join(rdd2)
joined.collect()
// Output: (1,(Alice,Math))
```

- ♦ **leftOuterJoin**

```
val left = rdd1.leftOuterJoin(rdd2)
left.collect()
// Output: (1,(Alice,Some(Math))), (2,(Bob,None))
```

- ♦ **rightOuterJoin**

```
val right = rdd1.rightOuterJoin(rdd2)
right.collect()
// Output: (1,(Some(Alice),Math)), (3,(None,Science))
```

- ♦ **fullOuterJoin**

```
val full = rdd1.fullOuterJoin(rdd2)
full.collect()
// Output: (1,(Some(Alice),Some(Math))), (2,(Some(Bob),None)),
(3,(None,Some(Science)))
```

- ♦ **Important: Joins require partitioning and shuffling.**

To optimize:

```
rdd1.partitionBy(new HashPartitioner(2)).join(rdd2)
```

Aggregations in Spark

1. RDD Aggregations

- ♦ `reduceByKey`

```
val scores = sc.parallelize(Seq(("math", 90), ("math", 80), ("eng", 85)))  
val totals = scores.reduceByKey(_ + _)
```

- ♦ `groupByKey`

```
val grouped = scores.groupByKey()
```

⚠ Not efficient for large shuffles. Prefer `reduceByKey` or `aggregateByKey`.

- ♦ `aggregateByKey`

```
val marks = sc.parallelize(Seq(("math", 90), ("math", 85), ("eng", 70)))  
val avg = marks.aggregateByKey((0, 0))(  
  (acc, value) => (acc._1 + value, acc._2 + 1),  
  (a, b) => (a._1 + b._1, a._2 + b._2)  
)  
val averages = avg.mapValues { case (sum, count) => sum / count }
```

Aggregations in Spark SQL / DataFrame API

```
import org.apache.spark.sql.functions._  
  
val df = Seq(  
  ("math", 90),  
  ("math", 85),  
  ("eng", 70)  
)  
  .toDF("subject", "score")
```

```
// Sum
df.groupBy("subject").agg(sum("score").as("total"))

// Average
df.groupBy("subject").agg(avg("score").as("average"))

// Count
df.groupBy("subject").agg(count("*").as("count"))

// Multiple Aggregations
df.groupBy("subject").agg(
  count("*").as("cnt"),
  sum("score").as("total"),
  avg("score").as("avg")
)
```

What Are Complex Transformations?

These are combinations or sequences of multiple transformations (e.g., map, filter, flatMap, join, groupByKey, reduceByKey, aggregateByKey, etc.) applied in a functional and efficient way to transform big datasets.

RDD-Based Complex Transformations

♦ 1. Chained Transformations Example

```
val rdd = sc.parallelize(Seq("spark is fun", "scala is powerful", "big data rocks"))
```

```
val words = rdd
```

```

.flatMap(_.split(" "))      // ["spark", "is", "fun", ...]
.filter(_.length > 2)       // remove short words
.map(word => (word.toLowerCase, 1))
.reduceByKey(_ + _)         // count frequencies

```

```
words.collect()
```

♦ 2. Nested Joins and Aggregations

```
val students = sc.parallelize(Seq((1, "Alice"), (2, "Bob"), (3, "Charlie")))
```

```
val scores = sc.parallelize(Seq((1, 80), (2, 90), (1, 70), (3, 85)))
```

```
val grouped = scores
```

```
.mapValues(score => (score, 1)) // (id, (score, 1))
```

```
.reduceByKey { case ((s1, c1), (s2, c2)) => (s1 + s2, c1 + c2) }
```

```
val avgScores = grouped.mapValues { case (sum, count) => sum / count
}
```

```
val result = students.join(avgScores)
```

```
result.collect()
```

```
// Output: (1, (Alice, 75)), (2, (Bob, 90)), (3, (Charlie, 85))
```

♦ 3. Multi-key Aggregation

```
val data = sc.parallelize(Seq(
```

```
  ("math", "Alice", 90),
```

```
  ("math", "Bob", 80),
```

```
  ("science", "Alice", 85),
```

```
  ("science", "Bob", 70)
```

```
))
```

```
val keyed = data.map { case (subject, student, marks) =>  
  ((subject, student), marks)  
}
```

```
val total = keyed.reduceByKey(_ + _)
```

```
total.collect()
```

```
// Output: ((math,Alice),90), ((math,Bob),80), ...
```

DataFrame-Based Complex Transformations

```
import spark.implicits._
```

```
import org.apache.spark.sql.functions._
```

```
val df = Seq(  
  ("Alice", "math", 90),  
  ("Alice", "science", 85),  
  ("Bob", "math", 80),  
  ("Bob", "science", 70)  
).toDF("name", "subject", "score")
```

```
// Pivot + Aggregation
```

```
val pivoted = df  
  .groupBy("name")  
  .pivot("subject")  
  .agg(avg("score"))
```

```
pivoted.show()
```

♦ 4. Window Functions (Advanced)

```
import org.apache.spark.sql.expressions.Window

val df = Seq(
  ("Alice", "2023-01-01", 100),
  ("Alice", "2023-01-02", 200),
  ("Bob", "2023-01-01", 50),
  ("Bob", "2023-01-02", 60)
).toDF("name", "date", "sales")

val windowSpec = Window.partitionBy("name").orderBy("date")

val withRunningTotal = df.withColumn("running_total",
  sum("sales").over(windowSpec))

withRunningTotal.show()
```

Optimization Techniques

1. Caching & Persistence

- Use `cache()` or `persist()` when reusing the same RDD/DataFrame multiple times.
- Avoid recomputation.

```
val df = spark.read.csv("big.csv").cache()
```

```
df.count()
```

```
df.groupBy("col1").count().show()
```

♦ Use `.persist(StorageLevel.MEMORY_AND_DISK)` if data is too large for memory.

2. Partitioning Strategy

- Repartition to increase parallelism:

```
val df2 = df.repartition(8, $"keyColumn")
```

- Coalesce to reduce partitions (less shuffling):

```
val df3 = df.coalesce(2)
```

♦ Rule of Thumb: Avoid too many small partitions or too few large ones.

3. Avoid Wide Transformations (Shuffles)

- `groupByKey`, `distinct`, `join`, etc., cause shuffles, which are expensive.
- Prefer `reduceByKey` over `groupByKey`.

```
// Better:
```

```
rdd.reduceByKey(_ + _)
```

```
// Avoid:
```

```
rdd.groupByKey().mapValues(_.sum)
```

4. Broadcast Joins

Use broadcast joins when joining a small dataset with a large one.

```
import org.apache.spark.sql.functions.broadcast
```

```
val result = largeDF.join(broadcast(smallDF), "id")
```

Reduces data shuffling.

5. Pushdown Predicates

Let Spark push filters to the source level (like Parquet, JDBC):

```
val df = spark.read
    .option("pushDownPredicate", true)
    .parquet("data.parquet")
    .filter($"status" === "active")
```

6. Column Pruning

Read only necessary columns:

```
val df = spark.read.parquet("data.parquet").select("id", "name")
```

Reduces I/O and memory usage.

7. Use Efficient File Formats

Use Parquet or ORC over CSV/JSON:

- Columnar
- Compressed
- Supports predicate pushdown & schema evolution

```
df.write.parquet("output/")
```

8. Optimize Joins

Join Type	Optimization Method
Large + Small	Broadcast join
Large + Large	Repartition on join key
Skewed Join	Salting / Skew join optimization

// Join optimization example

```
df1.repartition($"key").join(df2.repartition($"key"), "key")
```

9. Use explain() and Spark UI

Check logical & physical plans:

```
df.explain(true)
```

Use Spark UI:

- Check stages, tasks, memory usage
- Identify slow operations or data skew

10. Tungsten & Catalyst Optimization (built-in)

Spark 2.x+ includes:

- Tungsten: memory management & binary processing engine
- Catalyst: logical & physical query optimization

➡ Just write high-level transformations (don't micro-optimize unless needed).

11. Avoid UDFs Unless Necessary

- UDFs prevent Catalyst optimization.
- Use built-in Spark SQL functions instead:

```
// Instead of this:
```

```
val myUDF = udf((x: String) => x.toUpperCase)
```

```
df.withColumn("name", myUDF($"name"))
```

```
// Use this:
```

```
df.withColumn("name", upper($"name"))
```

12. Memory & Executor Tuning

Set via spark-submit or Spark config:

Config	Purpose
spark.executor.memory	Executor heap memory

Config	Purpose
spark.executor.cores	Number of cores per executor
spark.sql.shuffle.partitions	Number of shuffle partitions
spark.memory.fraction	Memory used for execution